




INTRODUCTION

- Convolutional Nearest Neighbor (ConvNN)** reinterprets convolution as k-nearest neighbor aggregation with flexible neighbor selection criteria.
- Standard convolution implicitly performs k-NN with fixed spatial distance (e.g., 3x3 kernel = k = 9 spatially-adjacent neighbors including self).
- ConvNN generalizes this by allowing neighbor selection based on:
 - Spatial distance (reduces to standard convolution)
 - Feature similarity (cosine/Euclidean)
 - Hybrid spatial-feature metrics
- Core Algorithm of ConvNN:
 - Compute pairwise similarities between all spatial positions
 - Select k-nearest neighbors per position via hard top-k
 - Aggregate neighbors with learnable weights (1D convolution)

BASE ALGORITHM

ConvNN Visualization



1. Similarity Computation

$$S = XX^T \in \mathbb{R}^{n \times n} \text{ where } S_{ij} = \mathbf{x}_i^T \mathbf{x}_j$$

2. K-Nearest Neighbor Selection

$$I_k = k - \text{argmax}(XX^T) \in \mathbb{R}^{n \times n}$$

$$\text{Neighbors} = X[I_k[i, :], :] \in \mathbb{R}^{k \times n}$$

Algorithm 1 Convolutional Nearest Neighbors 1D


Input: $\mathbf{X} \in \mathbb{R}^{B \times C \times N}$ (batch \times channels \times tokens)

Parameters: k (number of neighbors)

Output: $\mathbf{Y} \in \mathbb{R}^{B \times C' \times N}$

```
1: // For each batch element
2: Let  $X = \mathbf{X}[b, :, :]^T \in \mathbb{R}^{N \times C}$  with columns  $\mathcal{X} = \{\mathbf{x}_i\}_{i=1}^N$ 
3:
4: // Step 1: Compute similarity matrix
5: Assume each  $\mathbf{x}_i$  is  $\ell_2$ -normalized:  $\|\mathbf{x}_i\|_2 = 1$ 
6: Compute similarity:  $S = XX^T \in \mathbb{R}^{N \times N}$  where  $S_{ij} = \mathbf{x}_i^T \mathbf{x}_j$ 
7:
8: // Step 2: Find k-nearest neighbors
9:  $I_k = \text{argmax}_k(S) \in \{0, 1\}^{N \times N}$ 
10:
11: // Step 3: Gather features
12: for  $i \in [1, N]$  do
13:    $\mathcal{N}_k(\mathbf{x}_i) = X[I_k[i, :], :] \in \mathbb{R}^{k \times C}$ 
14:    $\mathbf{V}_{\text{prime}}[:, i, i \cdot k : (i+1) \cdot k] = \mathcal{N}_k(\mathbf{x}_i)^T$ 
15: end for
16:
17: // Step 4: Convolve
18:  $\mathbf{Y} = \text{Conv1D}(\mathbf{V}_{\text{prime}}, \text{kernel\_size} = k, \text{stride} = k)$ 
19:
20: return  $\mathbf{Y}$ 
```

SIMILARITY COMPUTATION SPEED-UPS



- To reduce $O(N^2)$ complexity of all to all similarity computation, we introduce two sampling methods: Random Sparsification and Spatial Sparsification.
- Trade-off between computational efficiency and neighbor selection quality is controlled by sampling parameter n .


ARCHITECTURE AND TRAINING

- Architecture:** VGG-11 with Conv2d layers replaced by ConvNN and branching layers
- Dataset:** CIFAR-10 image classification
- Training:** 60 epochs with AdamW (lr=1e-5, wd=1e-6), StepLR scheduler (gamma=0.95, step=2)
- Variants tested:**
 - Location-only (spatial distance)
 - Feature-only (cosine similarity)
 - Hybrid (weighted combination)
 - Branching with ratio (e.g., 50% Conv2d + 50% ConvNN)

RESULTS

Training and Test Loss


Comparison of Conv2d and Branching ConvNN



Branching ConvNN = Branching with branching ratio 0.500, kernel_size = 3, K = 9, Feature Similarity and Aggregation.

Model Accuracy by Kernel Size and Type

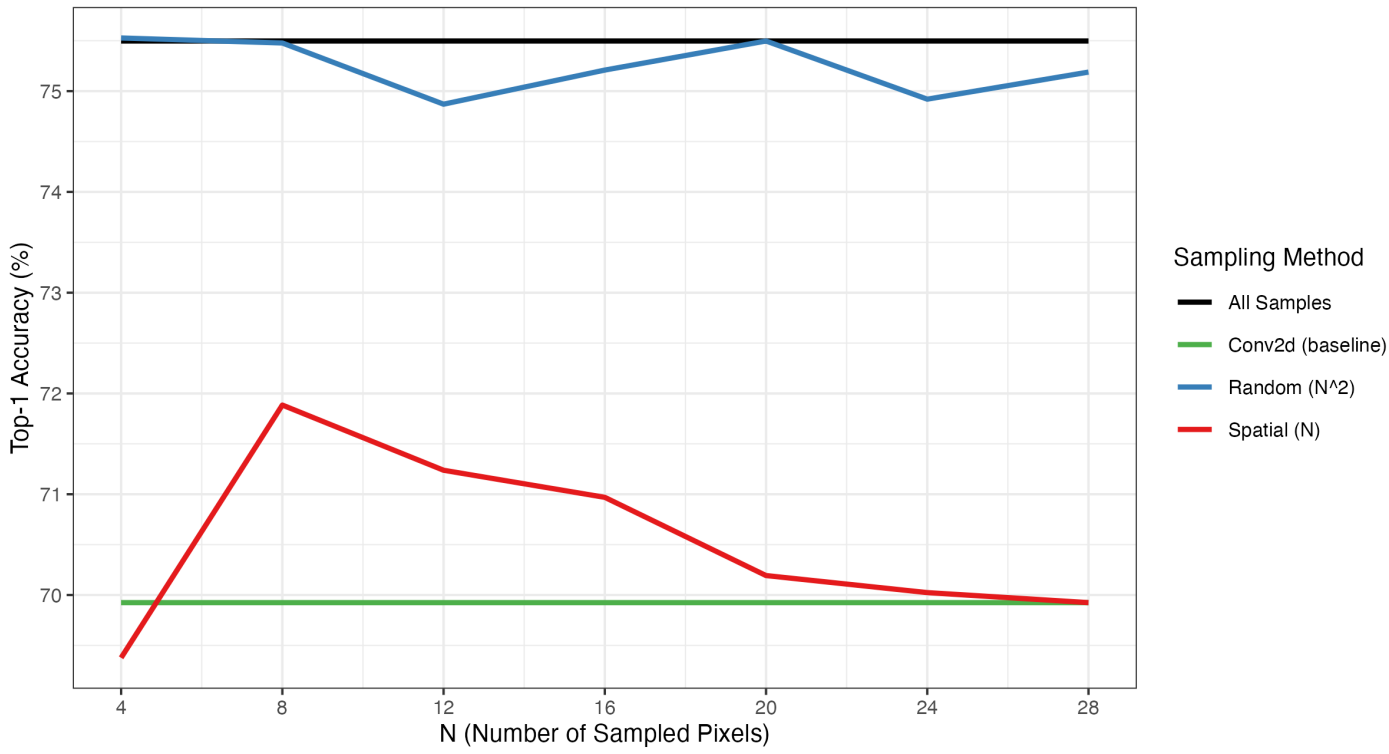
Comparison of Conv2d and Branching ConvNN



Branching ConvNN = Branching with branching ratio 0.250, Location + Feature Similarity and Aggregation.

Model Performance vs. N

Top-1 Accuracy for Random and Spatial Sampling Methods



Branching ConvNN = Branching with branching ratio 0.250, Location + Feature Similarity and Aggregation. Spatial Sampling = $N = N \times N$ sub grid 3, Random Sampling = N^2 pixels.

CONVOLUTION AND ATTENTION

1. Convolution

$$S = D = 2(1 - X^T X) \in \mathbb{R}^{n \times n} \text{ where } D_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2^2 = 2(1 - \mathbf{x}_i^T \mathbf{x}_j)$$

$$I_k = k - \text{argmax}(2(1 - X^T X)) \in \mathbb{R}^{n \times n}$$

$$\text{Neighbors} = X[I_k[i, :], :] \in \mathbb{R}^{k \times n}$$

2. Convolutional Nearest Neighbor

$$S = XX^T \in \mathbb{R}^{n \times n} \text{ where } S_{ij} = \mathbf{x}_i^T \mathbf{x}_j$$

$$I_k = k - \text{argmax}(XX^T) \in \mathbb{R}^{n \times n}$$

$$\text{Neighbors} = X[I_k[i, :], :] \in \mathbb{R}^{k \times n}$$

3. Attention

$$QK^T \in \mathbb{R}^{n \times n} \text{ where } Q = w_Q X, K = w_K X$$

$$A(Q, K) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \in \mathbb{R}^{n \times n}$$

$$\text{Attention}(Q, K, V) = A(Q, K)V \text{ where } V = w_V X$$

DISCUSSION

- Hybrid similarity** (spatial + feature) outperforms pure spatial or pure feature selection
- Branching architecture** achieves best performance by combining ConvNN's global context with Conv2d's spatial locality.
- ConvNN unifies convolution and attention as neighbor aggregation differ:
 - Spatial-only \rightarrow standard convolution
 - All positions with soft weights with linear projection \rightarrow self-attention
 - ConvNN occupies the middle ground with hard, content-aware selection
- Feature work:** Extend to Vision Transformers, explore learnable similarity metrics, investigate soft vs. hard selection.

REFERENCES

- A. Buades, B. Coll and J. -M. Morel, "A non-local algorithm for image denoising," *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, San Diego, CA, USA, 2005, pp. 60-65 vol. 2, doi: 10.1109/CVPR.2005.38.
- Singh, Sidak Pal, and Martin Jaggi. "Model fusion via optimal transport." *Advances in Neural Information Processing Systems* 33 (2020): 22045-22055.
- Plötz, Tobias, and Stefan Roth. "Neural nearest neighbors networks." *Advances in Neural information processing systems* 31 (2018).
- Wang, Xiaolong, et al. "Non-local neural networks." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2018.
- Vaswani, Ashish, et al. "Attention is all you need." *Advances in neural information processing systems* 30 (2017).